

Untitled(3)

February 26, 2019

```
In [4]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from scipy.special import xlogy
import pickle
```

0.1 Class Layer:

- Each layer has many features input_dim, output_dim, activation_function, its derivative, matrix W (weights) and b (biases).
- Each layer also has other parameters like A(input to layer), Z(output from layer), dW(gradients change in W), db(gradients change in b). However, these variables are temporary and removed as soon as their work is done.

0.2 Class Neural Network:

- Variables list of layers, Number of iterations, learning rate(alpha), and batch size)
- It also contain all the activation functions as well as their derivatives.
- It contain functions for fitting and predicting data.
- There are functions for computing loss, computing error, computing accuracy.
- There are functions which are used while training data like forward_propagation, backward_propagation, update_parameters.

Forward Propagation: For each layer there is input $A[l]$ to it. Each layer has $W[l]$ and $b[l]$ pre-computed. Moreover there is an activation function $g[l]$ already associated to each layer. Two values $Z[l]$ and $A[l+1]$

$$Z^l = W^l \cdot A^l + b^l$$
$$A^{l+1} = g^{l+1}(Z^l)$$

Then this A^{l+1} is forwarded as input to next layer. then final output is computed.

Cost function: Cost is computed as Cross entropy which is simply summation over original value and log of predicted value i.e.

$$L\{Y, A\} = \frac{-1}{m} \sum_{k=1}^m (Y * \log(A) + ((1 - Y) * \log(1 - A)))$$

it derivative is given as:

$$dA\{Y, A\} = \frac{-Y}{A} + \frac{1 - Y}{1 - A}$$

Backward propogation For each layer **dAl** is given as input from it **dZl** is computed. Which in then used to compute **dWl** and **dbl**. We also compute **dAl-1**, it is passed as error for the previous layer.

$$dZ^l = dA^l * g'(Z^l)$$

$$dA^{l-1} = W^l . dZ^l$$

$$dW^l = \frac{1}{m} * dZ^l . A^l$$

$$db^l = \frac{1}{m} * \sum_{k=1}^m dZ^l$$

$$W^l = W^l - \alpha * dW^l$$

$$b^l = b^l - \alpha * db^l$$

```
In [20]: class layer:
    def __init__(self, output_dim, input_dim, activation_function, derivative):
        self.output_dim = output_dim
        self.input_dim = input_dim
        self.activation_function = activation_function
        self.activation_function_derivative = derivative
        self.W = np.random.randn(output_dim, input_dim)*np.sqrt(2/input_dim)
        self.b = np.zeros((output_dim, 1))

    def print_layer_detail(self):
        print(
            """Input_dim = {0}, output_dim = {1}, activation_function = {2}, W.shape =
            Z.shape = {4}, b.shape = {5}""".format(
                self.input_dim, self.output_dim,
                self.activation_function.__name__, self.W.shape, self.Z.shape,
                self.b.shape))

class NeuralNetwork:
    def __init__(self, iterations=100, alpha=0.01, batch_size=50):
        self.layers = []
        self.iterations = iterations
        self.alpha = alpha
        self.epsilon = 1e-11
```

```

        self.batch_size = batch_size

def change_iterations(self, iterations):
    self.iterations = iterations

def change_alpha(self, alpha):
    self.alpha = alpha

def sigmoid(self, z):
    s = 1 / (1 + np.exp(-z))
    return s

def sigmoid_derivative(self, z):
    s = self.sigmoid(z)
    s = s*(1-s)
    return s

def relu(self, z):
    #     print(z)
    s = np.maximum(0, z)
    return s

def relu_derivative(self, z):
    s = (z>0)
    #z[z >= 0] = 1
    #z[z < 0] = 0
    s = s.astype('int')
    assert(s.shape == z.shape)
    return s

def tanh(self, z):
    return np.tanh(z)

def tanh_derivative(self, z):
    return (1-np.square(np.tanh(z)))

def softmax(self, z):
    z = z - z.max(axis=0, keepdims=True)
    y = np.exp(z)
    y = np.nan_to_num(y)
    y = y / y.sum(axis=0, keepdims=True)
    return y

def softmax_derivative(self, z):
    print(z.shape)
    z = z - z.max(axis=0, keepdims=True)
    y = np.exp(z)
    y = (y * (y.sum(axis=0, keepdims=True) - y)) / np.square(

```

```

        y.sum(axis=0, keepdims=True))
    return y

def standardize(self, X):
    X_standardized = (X-self.mean)/self.std
    #         display(X_standardized, X_standardized.shape)
    return X_standardized

def add_layer(self, output_dim, input_dim, activation="relu"):
    activation_function = None
    derivative = None
    if activation == "relu":
        activation_function = self.relu
        derivative = self.relu_derivative
    elif activation == "sigmoid":
        activation_function = self.sigmoid
        derivative = self.sigmoid_derivative
    elif activation == "tanh":
        activation_function = self.tanh
        derivative = self.tanh_derivative
    elif activation == "softmax":
        activation_function = self.softmax
        derivative = self.softmax_derivative
    else:
        raise ("Not a valid error function")
        return
    new_layer = layer(output_dim, input_dim, activation_function,
                      derivative)
    if len(self.layers) == 0:
        self.input_shape = input_dim
    self.output_shape = output_dim
    self.layers.append(new_layer)

def fit(self, X, y):
    self.mean, self.std = X.mean(), X.std()
    X = self.standardize(X)
    X = np.array(X).T
    y = np.array(y).T
    assert (X.shape[0] == self.input_shape)
    assert (y.shape[0] == self.output_shape)
    assert (X.shape[1] == y.shape[1])
    m = X.shape[1]
    costs = []

    #### X.shape = (number_of_features, number_of_rows)
    #### y.shape = (number_of_labels, number_of_rows)
    full_X = X.T
    full_y = y.T

```

```

for i in range(self.iterations):
    p = np.random.permutation(m)
    full_X, full_y = full_X[p], full_y[p]
    print("Iteration Number: ", i+1)
    start = 0
    end = self.batch_size
    xxx = 2*m/(self.batch_size*3)
    while end <= m:
        X = full_X.T[:, start:end]
        y = full_y.T[:, start:end]
        start+=self.batch_size
        end+=self.batch_size
        if end%xxx == 0:
            print("#", end='')
        #### Forward Propagation
        A = X
        for layer_no in range(len(self.layers)):
            A = self.forward_propagation(layer_no, A)
            A = np.nan_to_num(A)
            A = A / A.sum(axis=0, keepdims=True)
            print(max(A.sum(axis = 1)))
            #### A.shape = (number_of_labels, number_of_rows)
        dZ = A - y
        W = self.layers[-1].W
        A = self.layers[-1].A
        dW = np.dot(dZ, A.T) / m
        #### shape of db = (output_dim, 1)
        db = (1 / m) * np.sum(dZ, axis=1, keepdims=True)
        #### shape of da_new = ((input_dim, output_dim)*(output_dim, number_of_labels))
        #### shape of da_new = (input_dim, number_of_rows)
        dA = np.dot(W.T, dZ)
        # print("da_new.shape = {}".format(da_new.shape))
        self.layers[-1].dW = dW
        self.layers[-1].db = db
        #### Backward Propagation
        for layer_no in range(len(self.layers) - 2, -1, -1):
            dA = self.backward_propagation(layer_no, dA)

        #### Update parameters

        for layer_no in range(len(self.layers)):
            self.update_parameters(i, layer_no)

    A = full_X.T
    for layer_no in range(len(self.layers)):
        A = self.forward_propagation(layer_no, A)
    cost = self.compute_cost(A, full_y.T)
    costs.append(cost)

```

```

#             print()

iteration_list = [i for i in range(1, len(costs) + 1)]
plt.style.use('fivethirtyeight')
plt.plot(iteration_list, costs)
plt.ylabel("Loss")
plt.xlabel("Iterations")
plt.show()
#         display(costs)

def forward_propagation(self, layer_no, A):
#         display("Layer number {}".format(layer_no))
    self.layers[layer_no].A = A
    ##### shape of A = (input_dim, number of rows)
    ##### shape of W = (output_dim, input_dim)
    W = self.layers[layer_no].W
    b = self.layers[layer_no].b
    ##### shape of b = (output_dim, 1)
    g = self.layers[layer_no].activation_function
    self.layers[layer_no].Z = np.dot(W, A) + b
    ##### shape of Z = (output_dim, number_of_rows)
    A = g(self.layers[layer_no].Z)
    ##### shape of A = (output_dim, number_of_rows)
    return A

def compute_cost(self, prediction, target):
#         display(y_hat.max())
# shape of prediction (number_of_labels, number of training rows)
    m = prediction.shape[1]
    clipped = np.clip(prediction, self.epsilon, 1 - self.epsilon)
    cost = target * np.log(clipped) + (1 - target) * np.log(1 - clipped)
    return -np.sum(cost)/m

def compute_error(self, prediction, target):
    denominator = np.maximum(prediction - prediction ** 2, self.epsilon)
    delta = (prediction - target) / denominator
#         delta = -np.nan_to_num(np.divide(target, prediction)) + np.nan_to_num(np.divide
    assert (delta.shape == target.shape == prediction.shape)
    return delta

def backward_propagation(self, layer_no, dA):
    ##### shape of dA = (output_dim, number_of_rows)
    ##### shape of W = (output_dim, input_dim)
    ##### shape of Z = (output_dim, number_of_rows)
    ##### shape of A = (input_dim, number_of_rows)
    W = self.layers[layer_no].W
    g_der = self.layers[layer_no].activation_function_derivative
    Z = self.layers[layer_no].Z

```

```

A = self.layers[layer_no].A
m = A.shape[1]
#### shape of dZ = (output_dim, number_of_rows)
#         print(
#             "dA.shape = {0}, Z.shape = {1}, activation_function = {2}".format
#             dA.shape, Z.shape, g_der.__name__)
dZ = dA * g_der(Z)
#### shape of dW = ((output_dim, number_of_rows)*(number_of_rows, input_dim))
#### shape of dW = (output_dim, input_dim)
dW = np.dot(dZ, A.T) / m
#### shape of db = (output_dim, 1)
db = (1 / m) * np.sum(dZ, axis=1, keepdims=True)
#### shape of da_new = ((input_dim, output_dim)*(output_dim, number_of_rows))
#### shape of da_new = (input_dim, number_of_rows)
da_new = np.dot(W.T, dZ)
#         print("da_new.shape = {0}".format(da_new.shape))
self.layers[layer_no].dW = dW
self.layers[layer_no].db = db
self.layers[layer_no].A = self.layers[layer_no].Z = None
return da_new

def update_parameters(self, iteration, layer_no):
    #### shape of W, dW = (output_dim, input_dim)
    #### shape of b, db = (output_dim, 1)
    W = self.layers[layer_no].W
    b = self.layers[layer_no].b
    dW = self.layers[layer_no].dW
    db = self.layers[layer_no].db
    alph = self.alpha/(1+iteration)
    W = W - alph * dW
    b = b - alph * db
    self.layers[layer_no].W = W
    self.layers[layer_no].b = b
    self.layers[layer_no].dW = self.layers[layer_no].db = None
#         display("W", W)
#         display("b", b)

def predict(self, X):
    X = np.array(self.standardize(X))
    X = np.array(X).T
    A = X
    for layer_no in range(len(self.layers)):
        A = self.forward_propagation(layer_no, A)
    A = A.T
#         for i in range(len(A)):
#             display(A[i])
    return A

def calculate_accuracy(self, y_pred, y_test):

```

```

y_hat = np.argmax(y_pred, axis=1).flatten()
y__test = np.array(y_test)
y__test = np.argmax(y__test, axis=1).flatten()
# print(np.unique(y_hat, return_counts=True))
count = 0
for yh, y in zip(y_hat, y__test):
    if (yh == y):
        count += 1
return (count / len(y_hat))

```

```

In [19]: data = pd.read_csv("Apparel/apparel-trainval.csv")
        data = data.astype('float64')

```

```

In [7]: display(data.head())

```

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	\
0	2.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	9.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	6.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	5.0	
3	0.0	0.0	0.0	0.0	1.0	2.0	0.0	0.0	0.0	
4	3.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

	pixel9	...	pixel775	pixel776	pixel777	pixel778	pixel779	pixel780	\
0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	...	0.0	0.0	0.0	30.0	43.0	0.0	
3	0.0	...	3.0	0.0	0.0	0.0	0.0	1.0	
4	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	

	pixel781	pixel782	pixel783	pixel784
0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0

[5 rows x 785 columns]

```

In [8]: y = pd.get_dummies(data['label'])

```

```

In [9]: X = data.drop('label', axis=1)

```

```

In [10]: display(X.head())
        display(y.head())

```

	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	\
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	5.0	0.0
3	0.0	0.0	0.0	1.0	2.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

	pixel10	...	pixel775	pixel776	pixel777	pixel778	pixel779	pixel780	\
0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	...	0.0	0.0	0.0	30.0	43.0	0.0	
3	0.0	...	3.0	0.0	0.0	0.0	0.0	1.0	
4	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	

	pixel781	pixel782	pixel783	pixel784
0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0

[5 rows x 784 columns]

	0.0	1.0	2.0	3.0	4.0	5.0	6.0	7.0	8.0	9.0
0	0	0	1	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	1
2	0	0	0	0	0	0	1	0	0	0
3	1	0	0	0	0	0	0	0	0	0
4	0	0	0	1	0	0	0	0	0	0

```
In [11]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42)
```


```
In [21]: model_tanh = NeuralNetwork(iterations=20, alpha=0.05, batch_size=50)
```

```
In [22]: model_tanh.add_layer(256, 784, "tanh")
        model_tanh.add_layer(128, 256, "tanh")
        model_tanh.add_layer(64, 128, "tanh")
        # model.add_layer(50, 100, "sigmoid")
        model_tanh.add_layer(10, 64, "softmax")
```

```
In [23]: model_tanh.fit(X_train, y_train)
```

```
Iteration Number: 1
#####Iteration Number: 2
#####Iteration Number: 3
#####Iteration Number: 4
#####Iteration Number: 5
#####Iteration Number: 6
#####Iteration Number: 7
```

```
#####Iteration Number: 8
#####Iteration Number: 9
#####Iteration Number: 10
#####Iteration Number: 11
#####Iteration Number: 12
#####Iteration Number: 13
#####Iteration Number: 14
#####Iteration Number: 15
#####Iteration Number: 16
#####Iteration Number: 17
#####Iteration Number: 18
#####Iteration Number: 19
#####Iteration Number: 20
#####
```



```
In [24]: with open('tanh.pkl', 'wb') as output:
          pickle.dump(model_tanh, output, pickle.HIGHEST_PROTOCOL)

In [25]: y_pred = model_tanh.predict(X_test)
          print("Accuracy : ",model_tanh.calculate_accuracy(y_pred, y_test)*100)


Accuracy : 88.36666666666667

In [26]: model_sigmoid = NeuralNetwork(iterations=20, alpha=0.2, batch_size=50)
          model_sigmoid.add_layer(256, 784, "sigmoid")
```

```
model_sigmoid.add_layer(128, 256, "sigmoid")
model_sigmoid.add_layer(64, 128, "sigmoid")
# model.add_layer(50, 100, "sigmoid")
model_sigmoid.add_layer(10, 64, "softmax")
```

```
In [27]: model_sigmoid.fit(X_train, y_train)
```

```
Iteration Number: 1
#####Iteration Number: 2
#####Iteration Number: 3
#####Iteration Number: 4
#####Iteration Number: 5
#####Iteration Number: 6
#####Iteration Number: 7
#####Iteration Number: 8
#####Iteration Number: 9
#####Iteration Number: 10
#####Iteration Number: 11
#####Iteration Number: 12
#####Iteration Number: 13
#####Iteration Number: 14
#####Iteration Number: 15
#####Iteration Number: 16
#####Iteration Number: 17
#####Iteration Number: 18
#####Iteration Number: 19
#####Iteration Number: 20
#####
```



```
In [28]: with open('sigmoid.pkl', 'wb') as output:
        pickle.dump(model_sigmoid, output, pickle.HIGHEST_PROTOCOL)

In [29]: y_pred = model_sigmoid.predict(X_test)
        print("Accuracy : ",model_sigmoid.calculate_accuracy(y_pred, y_test)*100)


Accuracy : 84.38333333333333
```

```
In [30]: model_relu = NeuralNetwork(iterations=20, alpha=0.01, batch_size=50)
```

```
In [31]: model_relu.add_layer(256, 784, "relu")
        model_relu.add_layer(128, 256, "relu")
        model_relu.add_layer(64, 128, "relu")
        model_relu.add_layer(10, 64, "softmax")
```

```
In [32]: model_relu.fit(X_train, y_train)
```

```
Iteration Number: 1
#####Iteration Number: 2
#####Iteration Number: 3
#####Iteration Number: 4
#####Iteration Number: 5
#####Iteration Number: 6
#####Iteration Number: 7
#####Iteration Number: 8
#####Iteration Number: 9
#####Iteration Number: 10
#####Iteration Number: 11
#####Iteration Number: 12
#####Iteration Number: 13
#####Iteration Number: 14
#####Iteration Number: 15
#####Iteration Number: 16
#####Iteration Number: 17
#####Iteration Number: 18
#####Iteration Number: 19
#####Iteration Number: 20
#####
```



```
In [33]: with open('relu.pkl', 'wb') as output:
          pickle.dump(model_relu, output, pickle.HIGHEST_PROTOCOL)

In [34]: y_pred = model_relu.predict(X_test)
          print("Accuracy : ",model_relu.calculate_accuracy(y_pred, y_test)*100)

Accuracy : 87.26666666666667
```

0.3 Best performing Architecture for Neural Network

model_tanh.add_layer(10, 64, "softmax") - Best performing architecture has 3 hidden layer - All these layer have tanh as activation function - Loss function used is Cross Entropy loss - Dimensions of layer - - Hidden Layer 1 : 784 * 256 - - Hidden Layer 2 : 256 * 128 - - Hidden Layer 3 : 128 * 64 - - Output Layer : 64 * 10

0.4 Effect of various activation function in hidden layer:


Different layers more or less perform decent. However, tanh seems to outperform others. Sigmoid is a slow learner function so with less number of iteration it is predicting comparatively less accurately. However the fastest convergence is observed in Relu. ##### Results on validation data: - Tanh 88.11% - sigmoid 83.75% - Relu 86.90%

0.5 Effect on Loss and Error of validation data with increasing number of layers


- With increasing number of layers the error and loss decrease in general. Thus model is able to fit the data.
- However after large number of layers it is increasing because of overfitting.
- Error function is smooth as compare to loss function.

```
In [21]: loss = []
         error = []
         for n_layers in range(1,11,1):
             model = NeuralNetwork(iterations=20, alpha=0.01, batch_size=50)
             model.add_layer(32, 784, "tanh")
             for i in range(n_layers):
                 model.add_layer(32, 32, "tanh")
             model.add_layer(10, 32, "softmax")
             model.fit(X_train, y_train)
             y_pred = model.predict(X_test)
             loss.append(model.compute_cost(y_pred.T, (y_test.values).T))
             # print(loss)
             error.append(1.0 - model.calculate_accuracy(y_pred, y_test))

In [22]: plt.plot([i for i in range(1,11,1)], error, label="Error")
         plt.legend(loc="best")
         plt.ylabel("Error")
         plt.xlabel("No. of layers")
         plt.show()
```



```
In [23]: plt.plot([i for i in range(1,11,1)], loss, label = "Cost")
plt.legend(loc = "best")
plt.ylabel("Error")
plt.xlabel("No. of layers")
plt.show()
```




0.6 Error on validation data with increasing Epochs

- With increasing epochs the error is decreasing.
- After 50 iterations the error seems to become constant.
- It is because after certain iterations the model is not learning new things.

```
In [24]: model = NeuralNetwork(iterations=1, alpha=0.05, batch_size=50)
model.add_layer(256, 784, "tanh")
model.add_layer(128, 256, "tanh")
model.add_layer(64, 128, "tanh")
# model.add_layer(50, 100, "sigmoid")
model.add_layer(10, 64, "softmax")
alpha = 0.05
error = []
for iteration in range(100):
    model.change_alpha(alpha/(1+iteration))
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    error.append(1.0 - model.calculate_accuracy(y_pred, y_test))
```

```
In [25]: plt.plot([i for i in range(100)], error, label = "Error")
plt.legend(loc = "best")
plt.ylabel("Error")
plt.xlabel("No. of layers")
plt.show()
```



```
In [35]: TEST_DATA = pd.read_csv("apparel-test.csv")
y_tanh = np.argmax(model_tanh.predict(TEST_DATA), axis=1).flatten()
y_relu = np.argmax(model_relu.predict(TEST_DATA), axis=1).flatten()
y_sigmoid = np.argmax(model_sigmoid.predict(TEST_DATA), axis=1).flatten()
output = pd.DataFrame()
output['tanh'] = y_tanh
output['relu'] = y_relu
output['sigmoid'] = y_sigmoid
```

```
In [38]: np.savetxt("output.csv", np.int32(y_tanh))
```

```
In [42]: tanh_file = open('tanh.pkl', "rb")
model_loaded = pickle.load(tanh_file)
y_pred_loaded = model_loaded.predict(X_test)
print(model_loaded.calculate_accuracy(y_pred_loaded, y_test))
```

```
0.8836666666666667
```


1 Question 2:

- Since it is a regression problem, by keeping the output activation function as $f(x) = x$ will give a linear combination of the hidden layer nodes, hence giving linear regression with Neural networks.

The hidden layer activations can be anything, but it would be preferable to use the function $f(x) = x$ as it is more intuitive for a linear regression, i.e linear combination of input features.

However one thing to note is that price of house can't be negative so we can use RELU activation function at output layer as it is essentially a linear function of values > 0 .